

Sensorimotor Rhythm or Aggregate Power? An Interpretability Audit of EEGPT on Motor Imagery

Authors [placeholder]. Affiliation [placeholder]. Track: Machine Learning and Computer Paradigms for Brain Discovery.


A linear classifier on EEGPT's frozen embeddings decodes motor imagery about as well as feature engineered baselines. Yet the embedding geometry itself is dominated by aggregate, largely aperiodic signal power. A parameterized, spatially controlled analysis finds only a small and consistent trend toward the sensorimotor mu rhythm, not a strong rhythm specific representation.

Introduction. The EEG power spectrum separates into a periodic component, the oscillations such as the sensorimotor mu and beta rhythms that index specific function [7], and an aperiodic $1/f$ component whose slope tracks excitation and inhibition balance and is itself a clinical biomarker [5]. Recent audits report that EEG foundation models lean on the aperiodic background, and the broadband power that comes with it, more than on the specific oscillations that carry task meaning [1, 2, 3, 4]. A model that reaches accuracy through aggregate power can still be useful, but it is not grounded in the specific rhythm that clinical and BCI interpretation of motor imagery ideally rests on [7]. We extend this line of audits to EEGPT (NeurIPS 2024), a model absent from them and architecturally distinct: it adds a spatio temporal representation alignment objective to masked reconstruction [6], making it a useful test of whether that architectural change alters what the model encodes. Answering it carefully matters, because raw band power conflates periodic and aperiodic activity [9] and spatial averaging can inflate apparent rhythm specificity, so the conclusion depends on doing the spectral decomposition and the spatial control properly.

Methods. We audit the frozen EEGPT encoder on PhysioNet motor imagery, imagery versus rest, across 35 subjects, using the 512 dimensional embeddings as they are used downstream. We decode imagery versus rest, and we run representational similarity analysis (RSA) [8] comparing the geometry of the embedding space to per channel spectral references. To avoid the band power conflation, we parameterize each trial spectrum into aperiodic ($1/f$) and periodic components [5, 10] and build the oscillatory reference from aperiodic adjusted power per channel, for motor bands (mu, beta) and non motor control bands (theta, gamma). We correlate against the embedding geometry controlling for task and for the aperiodic component, contrast central against occipital channels on mu (a scalp-topography distinction, not source localization) to separate sensorimotor mu from posterior alpha, and test the mu effect against a trial shuffling null.

Results. A linear classifier on the frozen embeddings decodes imagery versus rest at about 71 percent within subject (chance 50), on par with feature engineered baselines, and transfers modestly across subjects. Its representational geometry is dominated by aggregate power: it correlates with total signal power at Spearman about 0.3, higher than with any single band. Once the aperiodic component is removed and the spatial construction is held fixed, only a small and consistent trend toward the sensorimotor mu rhythm remains. Aperiodic adjusted mu power is tracked about 0.04 above the non motor control bands (mu above theta p approximately 0.002, mu above gamma p approximately 0.0003, in roughly 80 percent of subjects), it is stronger over central than occipital channels and survives controlling for occipital mu (p approximately $1e-6$), a scalp-topographic distinction (not source localization) consistent with sensorimotor mu rather than posterior alpha, and it exceeds a trial shuffling null (p approximately 0.001). Beta shows no reliable trend. We read this as a trend, not an established component: at $n = 35$ it is statistically detectable but small, and whether it persists and how precisely it can be estimated as sample size grows is an open question, since significance at a given sample size reflects detectability rather than magnitude.

Significance. EEGPT, despite its added alignment objective, represents motor imagery mainly through aggregate, largely aperiodic power, with only a small mu trend on top. This extends the aperiodic and aggregate power bias reported for reconstruction based models to a hybrid architecture, and it carries a methodological caution for the audit literature itself: whether a foundation model looks rhythm specific depends on the analysis, because raw band power conflates periodic and aperiodic activity [9] and spatial averaging inflates apparent specificity. The balanced reading is that EEGPT reliably captures a coarse scale signal, aggregate and largely aperiodic power, which is itself physiologically meaningful since the $1/f$ exponent indexes excitation and inhibition balance and is a biomarker [5], but it has not yet shown strong specificity to the finer sensorimotor rhythm, despite an architecture meant to capture consistent structure. This coarse but not fine pattern may be a general property of current self supervised pretraining rather than specific to EEGPT: its objective already combines masked reconstruction with a JEPA-like latent alignment, the two families most common in EEG foundation models, yet fine specificity does not clearly emerge. That motivates applying the same rhythm versus power lens to other foundation models and to emerging architectures such as JEPA, rather than assuming a newer objective resolves the limitation. Whether the small mu trend strengthens with more subjects, and whether such objectives can be steered toward finer oscillatory structure, are open questions we are pursuing.



References

[1] Aperiodic and Low Frequency Spectral Bias in Reconstruction based EEG Foundation Models. arXiv:2605.26434, 2026. [2] Tang et al. What Do EEG Foundation Models Capture from Human Brain Signals? arXiv:2605.11410, 2026. [3] Beyond Accuracy: Robustness, Interpretability and Expressiveness of EEG Foundation Models. arXiv:2605.17562, 2026. [4] Shama, Amornsirikul, Venkataraman. EEG-PRISM: Physiologically-Grounded Interpretability of Predictions by EEG Foundation Models. arXiv:2608.13676, 2026. [5] Donoghue et al. Parameterizing neural power spectra into periodic and aperiodic components. *Nature Neuroscience*, 2020. [6] Wang et al. EEGPT: Pretrained Transformer for Universal and Reliable Representation of EEG Signals. *NeurIPS 2024*. [7] Pfurtscheller and Lopes da Silva. Event related EEG/MEG synchronization and desynchronization: basic principles. *Clinical Neurophysiology*, 1999. [8] Kriegeskorte, Mur, Bandettini. Representational similarity analysis. *Frontiers in Systems Neuroscience*, 2008. [9] Donoghue, Dominguez, Voytek. Electrophysiological frequency band ratio measures conflate periodic and aperiodic neural activity. *eNeuro*, 2020. [10] Gerster et al. Separating neural oscillations from aperiodic $1/f$ activity: challenges and recommendations. *Neuroinformatics*, 2022.